

```
In [46]: #Import Libraries
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from pandas.plotting import scatter_matrix
%matplotlib inline
```

```
In [47]: #read data and display the first 5 rows
```

```
auto_df = pd.read_excel("/content/Sales_Data.xlsx")
auto_df.head()
```

Out[47]:

	ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	ORDERDATE	DAY_SINCE_LASTORDER	STA
0	10107	30	95.70		2 2871.00	43155		828 Ship
1	10121	34	81.35		5 2765.90	43227		757 Ship
2	10134	41	94.74		2 3884.34	43282		703 Ship
3	10145	45	83.26		6 3746.70	43337		649 Ship
4	10168	36	96.66		1 3479.76	43401		586 Ship



```
In [48]: #View data information....data types for columns
auto_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2747 entries, 0 to 2746
Data columns (total 20 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   ORDERNUMBER      2747 non-null    int64  
 1   QUANTITYORDERED 2747 non-null    int64  
 2   PRICEEACH        2747 non-null    float64 
 3   ORDERLINENUMBER 2747 non-null    int64  
 4   SALES            2747 non-null    float64 
 5   ORDERDATE        2747 non-null    int64  
 6   DAYS_SINCE_LASTORDER 2747 non-null    int64  
 7   STATUS            2747 non-null    object  
 8   PRODUCTLINE      2747 non-null    object  
 9   MSRP              2747 non-null    int64  
 10  PRODUCTCODE      2747 non-null    object  
 11  CUSTOMERNAME     2747 non-null    object  
 12  PHONE             2747 non-null    object  
 13  ADDRESSLINE1     2747 non-null    object  
 14  CITY              2747 non-null    object  
 15  POSTALCODE       2747 non-null    object  
 16  COUNTRY           2747 non-null    object  
 17  CONTACTLASTNAME  2747 non-null    object  
 18  CONTACTFIRSTNAME 2747 non-null    object  
 19  DEALSIZE          2747 non-null    object  
dtypes: float64(2), int64(6), object(12)
memory usage: 429.3+ KB
```

The data constitutes of the following data types;

1. 6 columns with int64 data types
2. 2 columns with float64 data types
3. The remaining columns have object data type

```
In [49]: #display column names
auto_df.columns
```

```
Out[49]: Index(['ORDERNUMBER', 'QUANTITYORDERED', 'PRICEEACH', 'ORDERLINENUMBER',
   'SALES', 'ORDERDATE', 'DAYS_SINCE_LASTORDER', 'STATUS', 'PRODUCTLINE',
   'MSRP', 'PRODUCTCODE', 'CUSTOMERNAME', 'PHONE', 'ADDRESSLINE1', 'CITY',
   'POSTALCODE', 'COUNTRY', 'CONTACTLASTNAME', 'CONTACTFIRSTNAME',
   'DEALSIZE'],
  dtype='object')
```

```
In [50]: #to make EDA easier we drop some fields and convert date fields to date format
temp = ['PHONE', 'ADDRESSLINE1', 'CITY', 'POSTALCODE', 'CONTACTLASTNAME', 'CONTACTFIRSTNAME']
auto_df.drop(temp, axis=1, inplace=True)
```

```
In [51]: auto_df.head()
```

```
Out[51]:
```

	ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	ORDERDATE	DAYS_SINCE_LASTORDER	STA	
0	10107	30	95.70		2	2871.00	43155	828	Ship
1	10121	34	81.35		5	2765.90	43227	757	Ship
2	10134	41	94.74		2	3884.34	43282	703	Ship
3	10145	45	83.26		6	3746.70	43337	649	Ship
4	10168	36	96.66		1	3479.76	43401	586	Ship



```
In [52]: #count the number of NaN values in each
auto_df.isnull().sum()
```

Out[52]:


	0
ORDERNUMBER	0
QUANTITYORDERED	0
PRICEEACH	0
ORDERLINENUMBER	0
SALES	0
ORDERDATE	0
DAYS_SINCE_LASTORDER	0
STATUS	0
PRODUCTLINE	0
MSRP	0
PRODUCTCODE	0
CUSTOMERNAME	0
COUNTRY	0
DEALSIZE	0

dtype: int64

There are no null values in the data set

In [53]:

```
#generate box plots for numerical variables to see how the values are distributed
plt.rcParams['figure.figsize'] = [18,16]
auto_df.plot(kind="box", subplots=True, layout=(4,4), sharex=False, sharey=False)
plt.show()
```




Most of the values are distributed normally around the mean apart from QUANTITYORDERED, PRICEEACH and SALES which have outliers

```
In [54]: country_order_counts = auto_df['COUNTRY'].value_counts().reset_index()
country_order_counts.columns = ['Country', 'Order_Count']

plt.figure(figsize=(15, 7))
sns.barplot(x='Country', y='Order_Count', data=country_order_counts, palette='viridis', hue='Country', legend=False)
plt.title('Distribution of Orders Across Countries')
plt.xlabel('Country')
plt.ylabel('Number of Orders')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```

Distribution of Orders Across Countries



Most of the orders are coming from the USA followed Spain, France and Australia in that order

```
In [55]: #Generate a table showing how the orders are distributed across counties in percentages
country_order_counts['Percentage'] = (country_order_counts['Order_Count'] / country_order_counts['Order_Count'].sum())
display(country_order_counts.sort_values(by='Percentage', ascending=False))
```

	Country	Order_Count	Percentage
0	USA	928	33.782308
1	Spain	342	12.449945
2	France	314	11.430652
3	Australia	185	6.734620
4	UK	144	5.242082
5	Italy	113	4.113578
6	Finland	92	3.349108
7	Norway	85	3.094285
8	Singapore	79	2.875865
9	Canada	70	2.548234
10	Denmark	63	2.293411
11	Germany	62	2.257008
12	Sweden	57	2.074991
13	Austria	55	2.002184
14	Japan	52	1.892974
15	Belgium	33	1.201311
16	Switzerland	31	1.128504
17	Philippines	26	0.946487
18	Ireland	16	0.582454

Orders from the USA constitute of 33.78% of the orders placed. Ireland has the fewest orders constituting of 0.58%

```
In [56]: #show the value counts for unique values under 'STATUS', 'PRODUCTLINE' and 'DEALSIZE'
display(auto_df['STATUS'].value_counts())
```

count**STATUS**

Shipped	2541
Cancelled	60
Resolved	47
On Hold	44
In Process	41
Disputed	14

dtype: int64

```
In [57]: #generate a chart showing the percentage of each status
status_counts = auto_df['STATUS'].value_counts().reset_index()
status_counts.columns = ['STATUS', 'Count']

# Calculate percentages
total_count = status_counts['Count'].sum()
status_counts['Percentage'] = (status_counts['Count'] / total_count) * 100


plt.figure(figsize=(10, 6))
bars = plt.bar(status_counts['STATUS'], status_counts['Count'], color='teal')

# Set Labels and title
plt.title('Distribution of Order Status (with Percentages)')
plt.xlabel('Order Status')
plt.ylabel('Count')
plt.xticks(rotation=45, ha='right') # Rotate status names for better readability

# Annotate with Percentages
for i, bar in enumerate(bars):
    height = bar.get_height()
    percentage_value = status_counts['Percentage'].iloc[i]

    plt.text(bar.get_x() + bar.get_width() / 2.,
```

```
height + (total_count * 0.002), # Position text slightly above the bar, scaled to total  
f'{percentage_value:.1f}%',  
ha='center',  
va='bottom',  
fontsize=10)  
  
plt.tight_layout()  
plt.savefig('status_bar_chart_with_percentages.png')  
plt.show()
```



the shipped orders consist of 92,5 percent of all the orders

In [58]: `display(auto_df['PRODUCTLINE'].value_counts())`

PRODUCTLINE	count
Classic Cars	949
Vintage Cars	579
Motorcycles	313
Planes	304
Trucks and Buses	295
Ships	230
Trains	77

dtype: int64

In [59]: `#generate a chart showing the percentage of each product line
status_counts = auto_df['PRODUCTLINE'].value_counts().reset_index()
status_counts.columns = ['PRODUCTLINE', 'Count']

Calculate percentages
total_count = status_counts['Count'].sum()
status_counts['Percentage'] = (status_counts['Count'] / total_count) * 100

plt.figure(figsize=(10, 6))
bars = plt.bar(status_counts['PRODUCTLINE'], status_counts['Count'], color='teal')

Set Labels and title
plt.title('Distribution of Order Status (with Percentages)')
plt.xlabel('Order Status')
plt.ylabel('Count')
plt.xticks(rotation=45, ha='right') # Rotate status names for better readability


Annotate with Percentages
for i, bar in enumerate(bars):`

```
height = bar.get_height()
percentage_value = status_counts['Percentage'].iloc[i]

plt.text(bar.get_x() + bar.get_width() / 2.,
         height + (total_count * 0.002), # Position text slightly above the bar, scaled to total
         f'{percentage_value:.1f}%',
         ha='center',
         va='bottom',
         fontsize=10)

plt.tight_layout()
plt.savefig('status_bar_chart_with_percentages.png')
plt.show()
```

Distribution of Order Status (with Percentages)



Classic cars constitute of 34.5%, Vintage are 21.1% and Motocycles are 11.4 percent. The productline with least products is Trains at 2.8%.

```
In [60]: display(auto_df['DEALSIZE'].value_counts())
```

count**DEALSIZE****Medium** 1349**Small** 1246**Large** 152**dtype:** int64

```
In [61]: #generate a chart showing the percentage of each DEALSIZE
status_counts = auto_df['DEALSIZE'].value_counts().reset_index()
status_counts.columns = ['DEALSIZE', 'Count']

# Calculate percentages
total_count = status_counts['Count'].sum()
status_counts['Percentage'] = (status_counts['Count'] / total_count) * 100


plt.figure(figsize=(10, 6))
bars = plt.bar(status_counts['DEALSIZE'], status_counts['Count'], color='teal')

# Set labels and title
plt.title('Distribution of Order Status (with Percentages)')
plt.xlabel('Order Status')
plt.ylabel('Count')
plt.xticks(rotation=45, ha='right') # Rotate status names for better readability

# Annotate with Percentages
for i, bar in enumerate(bars):
    height = bar.get_height()
    percentage_value = status_counts['Percentage'].iloc[i]

    plt.text(bar.get_x() + bar.get_width() / 2.,
             height + (total_count * 0.002), # Position text slightly above the bar, scaled to total
             f'{percentage_value:.1f}%',
             ha='center',
             va='bottom',
             fontsize=10)
```

```
plt.tight_layout()  
plt.savefig('status_bar_chart_with_percentages.png')  
plt.show()
```




The Large orders are very few constituting of 5.5% while Medium and Small constitute 49.1% and 45.4% respectively.


```
In [62]: #convert the ORDERDATE to a date variable  
auto_df['ORDERDATE'] = pd.to_datetime(auto_df['ORDERDATE'])  
auto_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2747 entries, 0 to 2746
Data columns (total 14 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   ORDERNUMBER      2747 non-null    int64  
 1   QUANTITYORDERED 2747 non-null    int64  
 2   PRICEEACH        2747 non-null    float64 
 3   ORDERLINENUMBER 2747 non-null    int64  
 4   SALES            2747 non-null    float64 
 5   ORDERDATE        2747 non-null    datetime64[ns]
 6   DAYS_SINCE_LASTORDER 2747 non-null  int64  
 7   STATUS            2747 non-null    object  
 8   PRODUCTLINE       2747 non-null    object  
 9   MSRP              2747 non-null    int64  
 10  PRODUCTCODE       2747 non-null    object  
 11  CUSTOMERNAME     2747 non-null    object  
 12  COUNTRY           2747 non-null    object  
 13  DEALSIZE          2747 non-null    object  
dtypes: datetime64[ns](1), float64(2), int64(5), object(6)
memory usage: 300.6+ KB
```

```
In [63]: numeric_df = auto_df.select_dtypes(include=np.number)
corr = numeric_df.corr()


sns.heatmap(corr, xticklabels=corr.columns, yticklabels=corr.columns, annot=True, cmap='coolwarm')
plt.rcParams['figure.figsize'] = [12, 10]
plt.show()
```





1. There is a fair correlation (0.81) between PRICEEACH and SALES as well as PRICEEACH and MSRP (0.78)
2. SALES and MSRP have a correlation of 0.63
3. SALES and QUANTITYORDERED have a correlation of 0.55.

```
In [64]: #We generate a scatter matrix  
scatter_matrix(auto_df)  
plt.show()
```



In [64]: